


ZipMix: compact three-seed supervised fine-tuning screen

7 methods × 3 matched training seeds · 2,048 supervised answer-label tokens per run

○ Seed 0 □ Seed 1 △ Seed 2 ◆ Mean + 95% training-seed interval - - - Base checkpoint · evaluated once

SciQ · target accuracy
500 fixed test items

CommonsenseQA · retention accuracy
500 fixed validation items



Intervals: 95% Student-t across training seeds; evaluation items are held fixed. Intervals are not clipped.

Descriptive screening only. Three pairs cannot yield an exact two-sided sign-test p-value below 0.25.

This compact comparison does not establish broad generalization or state-of-the-art superiority.